

Audion Office OCR AI

[Русский](#) · [User Guide](#) · [History](#)

Contents

- [Why It Exists](#)
- [What Is in the Package](#)
- [What Is Assembled From It](#)
- [Principles](#)
- [Next](#)
- [Technical Reference](#)
 - [Workbench Naming](#)
 - [Second Pass](#)
 - [Full Archive](#)

Recognising scans and converting office documents: from paper to an editable document, a spreadsheet, a searchable PDF, and everything else — from a single recognition pass.

Why It Exists

Recognition costs money and time. The usual workflow ignores that: you run a scan, get a Word file — and a week later you need the spreadsheet. So you run it again. Then you need a PDF with a text layer — once more.

Every pass means new money, a new queue, and a new result that may differ from the previous one.

Here the order is different:

`sources → recognition → document package → review → any format`

You pay for recognition once. The result is stored in an independent `<name>.document` package, and every format afterwards is assembled from it locally — with no call to a paid service.

What Is in the Package

Not the text, but everything the text came from:

- an exact copy of the original;
- images of every page before and after preprocessing;
- hashes;
- the text and coordinates of **all** recognition candidates, not just the chosen ones;
- reading order, tables, merged cells;
- confidence for every fragment;
- cross-checks between different engines;
- the provider's raw response;
- room for manual corrections.

Completeness here means no loss of source data, not a promise of flawless recognition. A recognition error is visible and fixable precisely because everything it was based on lies beside it.

What Is Assembled From It

Eight results, independently and at once: DOCX, XLSX, PDF with a text layer, ODT, Markdown, recognition JSON, HTML, and a review sheet.

By default only DOCX is produced — there are no format presets, nothing superfluous is generated. Reassembly runs locally; no third-party office suite is required. The maximum page size is A3.

Principles

The second pass by another engine is independent. It receives the same cleaned raster but does not see the first engine's text. So a disagreement between them means genuine ambiguity, rather than a hint repeated back.

Tables are parsed by a physical grid. For ruled tables a grid is built from the long lines of the raster, and the column count is verified against a **1...N** numbering row taken from the coordinates of recognised words. Only then are the engines compared inside exact cells.

Merges are restored **only in the header** — there they are predictable. In the body of a table a merged cell more often means a parsing error than the author's intent.

Disagreements over digits are preserved. If the engines read a number differently, both readings stay in the package. A number is where a recognition error costs most, and a person should settle it.

Markdown is not the centre of the system. It is convenient as a control format and as input for language models, but the document package is not built around it.

Next

- [User Guide](#) — step by step.
 - [History](#) — what changed.
-

Technical Reference

Workbench Naming

The path rows are called **Source** and **Target**. The buttons are the same as in every Audion project: **Source**, **Add file...**, **Target**, **Reset**, **Delete**, **List**.

The source may be a single file or a folder. An external path is mirrored into the managed working folder, and results are synchronised into the chosen target.

Second Pass

setting	what it does
none	the main engine only
Tesseract	local cross-check
Yandex	cloud cross-check
Yandex and Tesseract	both

Full Archive

Packing the whole package into one archive remains an engine capability for compatibility, but it is not shown among the main buttons: it is rarely needed and takes a great deal of space.